

ANSIBLE

ANSIBLE BEST PRACTICES: THE ESSENTIALS

Ansible Automates: DC


Jamie Duncan
@jamieeduncan
cloudguy@redhat.com

about jduncan

6+ years with Red Hat



My daughter Elizabeth
#cutestThingEver



Coming Soon
#shamelessPlug

THIS SESSION IS ABOUT NUTS AND BOLTS

Roadmaps are great. This is not one of them.

For this session, I'm making the assumption that you're currently writing Ansible playbooks.

My goal is to help you make those playbooks more effective.

AUTOMATION == DOCUMENTATION

If done properly, the process of automating a process can **become** the documentation for the process.


Everything in Ansible revolves around this core concept.

Treat Ansible content like application code

Version control is your best friend

Start as simple as possible and iterate

- Start with a basic playbook and static inventory
- Refactor and modularize progressively as you and your environment mature



Do It with Style

- Create a style guide for all contributors
- Consistency in:
 - Tagging
 - Whitespace
 - Naming of Tasks, Plays, Variables, and Roles
 - Directory Layouts
- Enforce the style

```
basic-project/  
├── inventory  
│   ├── group_vars  
│   ├── host_vars  
│   └── hosts  
└── site.yml
```

```
myapp/
├── roles
│   ├── myapp
│   │   ├── tasks
│   │   │   └── main.yml
│   │   └── etc.etc
│   ├── nginx
│   │   └── etc.etc
│   └── proxy
│       └── etc.etc
└── site.yml
```



```
myapp/
├── config.yml
├── provision.yml
├── roles
│   └── requirements.yml
└── site.yml
```

Give inventory nodes human-meaningful names rather than IPs or DNS hostnames.

10.1.2.75

10.1.5.45

10.1.4.5

10.1.0.40



db1 ansible_host=10.1.2.75

db2 ansible_host=10.1.5.45

db3 ansible_host=10.1.4.5

db4 ansible_host=10.1.0.40

w14301.acme.com

w17802.acme.com

w19203.acme.com

w19304.acme.com

web1 ansible_host=w14301.acme.com

web2 ansible_host=w17802.acme.com

web3 ansible_host=w19203.acme.com

web4 ansible_host=w19203.acme.com

Group hosts for easier inventory selection and less conditional tasks -- the more groups the better.


WHAT	WHERE	WHEN
[db] db[1:4]	[east] db1 web1	[dev] db1 web1
[web] web[1:4]	db3 web3	[test] db3 web3
	[west] db2 web2 db4 web4	[prod] db2 web2 db4 web4

Use a single source of truth if you have it -- even if you have multiple sources, Ansible can unify them.

Stay in sync automatically

Reduce human error

Use your instance and provider metadata for more than pretty columns in your TPS reports



Proper variable naming can make plays more readable and avoid variable name conflicts

Use descriptive, unique human-meaningful variable names

Prefix variables with it's "owner" such as a role name, service, or package

```
apache_max_keepalive: 25
apache_port: 80
tomcat_port: 8080
```

Make the most of variables

Find the appropriate place for your variables based on what, where and when they are set or modified

Separate logic (tasks) from variables to reduce repetitive patterns and provided added flexibility.

```
- name: Clone student lesson app for a user
host: nodes
tasks:
  - name: Create ssh dir
    file:
      state: directory
      path: /home/{{ username }}/.ssh

  - name: Set Deployment Key
    copy:
      src: files/deploy_key
      dest: /home/{{ username }}/.ssh/id_rsa

  - name: Clone repo
    git:
      accept_hostkey: yes
      clone: yes
      dest: /home/{{ username }}/lightbulb
      key_file: /home/{{ username }}/.ssh/id_rsa
      repo: git@github.com:example/apprepo.git
```

EXHIBIT A

Embedded parameter values and repetitive home directory value pattern in multiple places

Works but could be more clearer and setup to be more flexible and maintainable

```
- name: Clone student lesson app for a user
host: nodes
vars:
  user_home: /home/{{ username }}
  user_ssh: "{{ user_home }}/.ssh"
  deploy_key: "{{ user_ssh }}/id_rsa"
  app_dest: "{{ user_home }}/exampleapp"
tasks:
  - name: Create ssh dir
    file:
      state: directory
      path: "{{ user_ssh }}"

  - name: Set Deployment Key
    copy:
      src: files/deploy_key
      dest: "{{ deploy_key }}"

  - name: Clone repo
    git:
      dest: "{{ app_dest }}"
      key_file: "{{ deploy_key }}"
      repo: git@github.com:example/exampleapp.git
```

EXHIBIT B

Parameters values are set thru values away from the task and can be overridden.

Human meaningful variables “document” what’s getting plugged into a task parameter

More easily refactored into a role

Use native YAML syntax to maximize the readability of your plays

- Vertical reading is easier
- Supports complex parameter values
- Works better with editor syntax highlighting in editors

NO!

- `name: install telegraf`
 `yum: name=telegraf-{{ telegraf_version }} state=present update_cache=yes disable_gpg_c`
 `notify: restart telegraf`
- `name: configure telegraf`
 `template: src=telegraf.conf.j2 dest=/etc/telegraf/telegraf.conf`
- `name: start telegraf`
 `service: name=telegraf state=started enabled=yes`

Better, but not quite all the way there...

- `name: install telegraf`
 `yum: >`
 `name=telegraf-{{ telegraf_version }}`
 `state=present`
 `update_cache=yes`
 `disable_gpg_check=yes`
 `enablerepo=telegraf`
 `notify: restart telegraf`
- `name: configure telegraf`
 `template: src=telegraf.conf.j2 dest=/etc/telegraf/telegraf.conf`
- `name: start telegraf`
 `service: name=telegraf state=started enabled=yes`

- name: install telegraf
yum:
 - name: telegraf-{{ telegraf_version }}
 - state: present
 - update_cache: yes
 - disable_gpg_check: yes
 - enablerepo: telegrafnotify: restart telegraf
- name: configure telegraf
template:
 - src: telegraf.conf.j2
 - dest: /etc/telegraf/telegraf.confnotify: restart telegraf
- name: start telegraf
service:
 - name: telegraf
 - state: started
 - enabled: yes



Names improve readability and user feedback

Give all your playbooks, tasks and blocks brief, reasonably unique and human-meaningful names

`$myvar` is never a good thing, and typing isn't that hard

EXHIBIT A

```
- hosts: web
  tasks:
    - yum:
        name: httpd
        state: latest

    - service:
        name: httpd
        state: started
        enabled: yes
```

```
PLAY [web]
*****
```

```
TASK [setup]
*****
ok: [web1]
```

```
TASK [yum]
*****
ok: [web1]
```

```
TASK [service]
*****
ok: [web1]
```

EXHIBIT B

```
- hosts: web
  name: installs and start apache
  tasks:
    - name: install apache packages
      yum:
        name: httpd
        state: latest

    - name: start apache service
      service:
        name: httpd
        state: started
        enabled: yes
```

```
PLAY [install and start apache]
*****
```


```
TASK [setup]
*****
ok: [web1]
```

```
TASK [install apache packages]
*****
ok: [web1]
```

```
TASK [start apache service]
*****
ok: [web1]
```

Focus avoids complexity

Keep plays and playbooks focused. Multiple simple playbooks are better than having a single, overburdened playbook full of conditional logic.




Clean up your debugging tasks

Make them optional with the verbosity parameter so they're only displayed when they are wanted.

- debug:
msg: "This always displays"
- debug:
msg: "This only displays with
ansible-playbook -vv+"
verbosity: 2

Don't just start services -- use smoke tests

```
- name: check for proper response
  uri:
    url: http://localhost/myapp
    return_content: yes
  register: result
  until: '"Hello World" in result.content'
  retries: 10
  delay: 1
```



Use command modules sparingly

- Use the run command modules like `shell` and `command` as a last resort
- Use the `command` module unless you really need I/O redirection that `shell` permits -- but be very careful.

Always seek out a module first

- name: add user
command: useradd appuser
- name: install apache
command: yum install httpd
- name: start apache
shell: |
service httpd start && chkconfig httpd on

- name: add user
user:
name: appuser
state: present
- name: install apache
yum:
name: httpd
state: latest
- name: start apache
service:
name: httpd
state: started
enabled: yes

Still using command modules a lot?

```
- hosts: all
  vars:
    cert_store: /etc/mycerts
    cert_name: my cert
  tasks:
    - name: check cert
      shell: certify --list --name={{ cert_name }} --cert_store={{ cert_store }} | grep "{{ cert_name }}"
      register: output

    - name: create cert
      command: certify --create --user=chris --name={{ cert_name }} --cert_store={{ cert_store }}
      when: output.stdout.find(cert_name) != -1
      register: output

    - name: sign cert
      command: certify --sign --name={{ cert_name }} --cert_store={{ cert_store }}
      when: output.stdout.find("created") != -1
```

Develop your own module! (seriously)

```
- hosts: all
vars:
  cert_store: /etc/mycerts
  cert_name: my cert
tasks:
- name: create and sign cert
  certify:
    state: present
    sign: yes
    user: chris
    name: "{{ cert_name }}"
    cert_store: "{{ cert_store }}"
```

Separate provisioning from deployment and configuration tasks

```
acme_corp/  
├── configure.yml  
├── provision.yml  
└── site.yml
```

```
$ cat site.yml
```

```
---
```

- import_playbook: provision.yml
- import_playbook: configure.yml

Jinja2 is powerful but you needn't use all of it

Templates should be simple:

- Variable substitution
- Conditionals
- Simple control structures/iterations
- Design your templates for your use case, not the world's
-


Things to avoid:

- Managing variables in a template
- Extensive and intricate conditionals
- Conditional logic based on embedded hostnames
- Complex nested iterations

Careful when mixing manual and automated configuration

Label template output files as being generated by Ansible

```
{{ ansible_managed | comment }}
```



Roles are the shareable unit of work in Ansible

- Like playbooks -- keep roles purpose and function focused
- Use a **roles/** subdirectory for roles developed for organizational clarity in a single project
- Follow the Ansible Galaxy pattern for roles that are to be shared beyond a single project
- Limit role dependencies

Sharing roles is paramount, and easy

- Use **ansible-galaxy init** to start your roles...
- ...then remove unneeded directories and stub files
- Use **ansible-galaxy** to install your roles -- even private ones
- Use a roles files (i.e. **requirements.yml**) to manifest any external roles your project is using
- Always specify a specific version such using a tag or commit for your roles

Command line tools have their limitations

- Coordination across a distributed teams & organization...
- Controlling access to credentials...
- Track, audit and report automation and management activity...
- Provide self-service or delegation...
- Integrate automation with enterprise systems...



Applications and Infrastructure are
continuously evolving.

Ansible is designed to do the same.

Thanks!



#AnsibleAutomates
@jamieeduncan